

HarvestGuard

This project focuses on Leaf Disease Classification using an Afrocentric dataset. The dataset is specifically curated to represent diverse agricultural scenarios across various African regions. It comprises annotated photos of leaves from a variety of crops, showcasing both healthy specimens and those affected by diseases. The goal is to develop a classification system leveraging machine learning techniques to accurately identify and categorize crop diseases based on the subtle symptoms observed in the annotated images. By using an Afrocentric dataset, the project aims to enhance the effectiveness of disease classification models tailored to the unique agricultural landscape of Africa (Responsible AI Lab, 2023)

```
# Import libraries
!pip install opendatasets
import pandas as pd
import cv2
import json
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import opendatasets as od
import glob
import os
import pathlib
import PIL
from PIL import Image
import PIL.Image
import PIL.ImageShow
from google.colab.patches import cv2_imshow
import numpy as np
from sklearn.preprocessing import MultiLabelBinarizer
from sklearn.model_selection import train_test_split
import pickle
from tensorflow import keras
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import opendatasets as od
import re
import pathlib
import imgaug as ia
ia.seed(1)
# imgaug uses matplotlib backend for displaying images
%matplotlib inline
from imgaug.augmentables.bbs import BoundingBox, BoundingBoxesOnImage
from keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.preprocessing.image import img_to_array
```

```

from imgaug import augmenters as iaa
# imageio library will be used for image input/output
import imageio
import PIL
from PIL import Image
import PIL.Image
import PIL.ImageShow
from google.colab.patches import cv2_imshow
import xml.etree.ElementTree as ET
from os import listdir
from sklearn.preprocessing import LabelBinarizer
from keras.layers import MaxPooling2D
from keras.layers import Activation, Flatten, Dropout, Dense
from keras import backend as K
from keras.preprocessing.image import ImageDataGenerator
from keras.optimizers.legacy import Adam
from keras.preprocessing import image
from tensorflow.keras.preprocessing.image import img_to_array
from keras.models import Sequential
from tensorflow.compat.v1.keras.layers import BatchNormalization
from keras.layers import Conv2D
import matplotlib.pyplot as plt
import xml.etree.ElementTree as ET
from tensorflow.keras.applications import VGG16
from tensorflow.keras.layers import Input
from tensorflow.keras.layers import Flatten
from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Input
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import img_to_array
from tensorflow.keras.preprocessing.image import load_img
from sklearn.utils import class_weight
import shutil
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
warnings.simplefilter(action='ignore', category=DeprecationWarning)

```

```

Requirement already satisfied: opendatasets in
/usr/local/lib/python3.10/dist-packages (0.1.22)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-
packages (from opendatasets) (4.66.1)
Requirement already satisfied: kaggle in
/usr/local/lib/python3.10/dist-packages (from opendatasets) (1.5.16)
Requirement already satisfied: click in
/usr/local/lib/python3.10/dist-packages (from opendatasets) (8.1.7)
Requirement already satisfied: six>=1.10 in
/usr/local/lib/python3.10/dist-packages (from kaggle->opendatasets)
(1.16.0)

```

```
Requirement already satisfied: certifi in
/usr/local/lib/python3.10/dist-packages (from kaggle->opendatasets)
(2023.11.17)
Requirement already satisfied: python-dateutil in
/usr/local/lib/python3.10/dist-packages (from kaggle->opendatasets)
(2.8.2)
Requirement already satisfied: requests in
/usr/local/lib/python3.10/dist-packages (from kaggle->opendatasets)
(2.31.0)
Requirement already satisfied: python-slugify in
/usr/local/lib/python3.10/dist-packages (from kaggle->opendatasets)
(8.0.1)
Requirement already satisfied: urllib3 in
/usr/local/lib/python3.10/dist-packages (from kaggle->opendatasets)
(2.0.7)
Requirement already satisfied: bleach in
/usr/local/lib/python3.10/dist-packages (from kaggle->opendatasets)
(6.1.0)
Requirement already satisfied: webencodings in
/usr/local/lib/python3.10/dist-packages (from bleach->kaggle-
>opendatasets) (0.5.1)
Requirement already satisfied: text-unidecode>=1.3 in
/usr/local/lib/python3.10/dist-packages (from python-slugify->kaggle-
>opendatasets) (1.3)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from requests->kaggle-
>opendatasets) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.10/dist-packages (from requests->kaggle-
>opendatasets) (3.6)
```

Afrocentric (African) Crop Dataset

The Afrocentric Crop Disease dataset on Kaggle. The dataset includes annotated images of crops from Ghana, categorizing them based on disease types. Each category contains information on disease categories, image sizes, bounding boxes, and sample images

```
# Download dataset from Kaggle to notebook directory
od.download('https://www.kaggle.com/datasets/responsibleailab/crop-
disease-ghana')

Skipping, found downloaded files in "./crop-disease-ghana" (use
force=True to force download)

# Print out test image
test_img =
'/content/crop-disease-ghana/input/Tomato/Tomato__Early_Blight/images/
20230518_134246.jpg'
```

```
test_img = cv2.imread(test_img)
#cv2_imshow(test_img)

def display_image(test_img):
    h, w = test_img.shape[0:2]
    neww = 300
    newh = int(neww*(h/w))
    test_img = cv2.resize(test_img, (neww, newh))
    cv2_imshow(test_img)
    cv2.waitKey(0)

display_image(test_img)
```



```
# Show files in directory
os.listdir('/content/crop-disease-ghana/input')
```

```
['Pepper',
'Tomato',
'Corn',
'dataset_labels.csv',
'label_map.pbtxt',
'category_index.pbtxt',
'label_map.txt',
'label_map.json']
```

```
# Read csv file with data labels
```

```
labels_df =
```

```
pd.read_csv('/content/crop-disease-ghana/input/dataset_labels.csv')
```

```
labels_df.head()
```

	filename	disease	crop	width	height
depth \					
0	20230524_104642.jpg	Corn Cercospora Leaf Spot	Corn	4080	1836
3					
1	20230524_104642.jpg	Corn Cercospora Leaf Spot	Corn	4080	1836
3					
2	20230524_104642.jpg	Corn Cercospora Leaf Spot	Corn	4080	1836
3					
3	20230524_104642.jpg	Corn Cercospora Leaf Spot	Corn	4080	1836
3					
4	20230524_104642.jpg	Corn Cercospora Leaf Spot	Corn	4080	1836
3					

	xmin	ymin	xmax	ymax	\
0	2052.653343	695.836619	2210.117161	785.809054	
1	1110.682288	901.086237	1228.780152	982.623756	
2	1647.746382	912.332791	1791.150930	1002.305226	
3	2491.302550	1275.034169	2589.717436	1367.818243	
4	3326.423156	1255.352699	3410.778773	1339.701857	

	ann_path	\
0	input\Corn\Corn__Cercospora_Leaf_Spot\annotati...	
1	input\Corn\Corn__Cercospora_Leaf_Spot\annotati...	
2	input\Corn\Corn__Cercospora_Leaf_Spot\annotati...	
3	input\Corn\Corn__Cercospora_Leaf_Spot\annotati...	
4	input\Corn\Corn__Cercospora_Leaf_Spot\annotati...	

	img_path
0	input\Corn\Corn__Cercospora_Leaf_Spot\images\2...
1	input\Corn\Corn__Cercospora_Leaf_Spot\images\2...
2	input\Corn\Corn__Cercospora_Leaf_Spot\images\2...
3	input\Corn\Corn__Cercospora_Leaf_Spot\images\2...
4	input\Corn\Corn__Cercospora_Leaf_Spot\images\2...

```
# Create new column combining width and height column
```

```
labels_df['shape'] = 0
```

```

shape_lst = []
for i in range(len(labels_df)):
    shape_lst.append((labels_df['width'][i], labels_df['height'][i]))

labels_df['shape'] = shape_lst
labels_df['shape'].value_counts()

(4032, 3024)      19548
(4080, 1836)      13500
(1920, 1280)      12418
(4080, 3060)       9007
(720, 480)        1473
(4000, 3000)       1123
(6720, 4480)        552
(1920, 2560)        428
(2560, 1920)        210
(2576, 1932)        102
Name: shape, dtype: int64

```

Addressing rare categories of diseases

```

# Merge similar classes
labels_df.replace(['Pepper Late Blight', 'Pepper Early Blight'], 'Pepper Leaf Blight', inplace = True)
labels_df.replace(['Tomato Late Blight', 'Tomato Early Blight'], 'Tomato Leaf Blight', inplace = True)

# Drop labels with too little data or multiple diseases
labels_df = labels_df[labels_df['disease'] != 'Tomato Mosaic']
labels_df = labels_df[labels_df['disease'] != 'Corn Northern Leaf Blight']
labels_df = labels_df[labels_df['disease'] != 'Pepper Leaf Mosaic']
labels_df['disease'].value_counts()

Corn Cercospora Leaf Spot      9431
Tomato Septoria                 9211
Tomato Leaf Blight              7993
Corn Streak                     4591
Tomato Healthy                  4066
Pepper Septoria                 3222
Pepper Bacterial Spot           2780
Corn Common Rust                2434
Corn Healthy                    2304
Pepper Leaf Curl                2175
Pepper Leaf Blight              1660
Tomato Fusarium                 1238
Pepper Healthy                  1049
Pepper Cercospora               704
Pepper Fusarium                 696
Tomato Leaf Curl                641

```

Tomato Bacterial Spot 639
Name: disease, dtype: int64

```
# Fix img path
loc_path = '/content/crop-disease-ghana/'
new_img = []
for txt in labels_df['img_path']:
    n = loc_path + txt.replace('\\', '/')
    new_img.append(n)
labels_df['img_path'] = new_img
labels_df['img_path'][0]

{"type": "string"}

# Fix annotation path
loc_path = '/content/crop-disease-ghana/'
new_ann = []
for txt in labels_df['ann_path']:
    a = loc_path + txt.replace('\\', '/')
    new_ann.append(a)
labels_df['ann_path'] = new_ann
labels_df['ann_path'][0]

{"type": "string"}

# Labels for categories
labels = {}
lab_len = len(np.unique(labels_df['disease']))
for j,i in zip(np.unique(labels_df['disease']), range(1, lab_len + 1)):
    labels[j] = i
print(labels)

{'Corn Cercospora Leaf Spot': 1, 'Corn Common Rust': 2, 'Corn Healthy': 3, 'Corn Streak': 4, 'Pepper Bacterial Spot': 5, 'Pepper Cercospora': 6, 'Pepper Fusarium': 7, 'Pepper Healthy': 8, 'Pepper Leaf Blight': 9, 'Pepper Leaf Curl': 10, 'Pepper Septoria': 11, 'Tomato Bacterial Spot': 12, 'Tomato Fusarium': 13, 'Tomato Healthy': 14, 'Tomato Leaf Blight': 15, 'Tomato Leaf Curl': 16, 'Tomato Septoria': 17}

# Use categorical labels to make new target column
enc_labels = []
for i in labels_df['disease']:
    for k,v in labels.items():
        if i == k :
            enc_labels.append(labels[i])
labels_df['label'] = enc_labels
labels_df.head()
```

	filename	disease	crop	width	height
depth \					


```

0 20230524_104642.jpg Corn Cercospora Leaf Spot Corn 4080 1836
3
1 20230524_104642.jpg Corn Cercospora Leaf Spot Corn 4080 1836
3
2 20230524_104642.jpg Corn Cercospora Leaf Spot Corn 4080 1836
3
3 20230524_104642.jpg Corn Cercospora Leaf Spot Corn 4080 1836
3
4 20230524_104642.jpg Corn Cercospora Leaf Spot Corn 4080 1836
3

```

```

      xmin      ymin      xmax      ymax \
0 2052.653343 695.836619 2210.117161 785.809054
1 1110.682288 901.086237 1228.780152 982.623756
2 1647.746382 912.332791 1791.150930 1002.305226
3 2491.302550 1275.034169 2589.717436 1367.818243
4 3326.423156 1255.352699 3410.778773 1339.701857

```

```

      ann_path \
0 /content/crop-disease-ghana/input/Corn/Corn__C...
1 /content/crop-disease-ghana/input/Corn/Corn__C...
2 /content/crop-disease-ghana/input/Corn/Corn__C...
3 /content/crop-disease-ghana/input/Corn/Corn__C...
4 /content/crop-disease-ghana/input/Corn/Corn__C...

```

```

      img_path      shape
label
0 /content/crop-disease-ghana/input/Corn/Corn__C... (4080, 1836)
1
1 /content/crop-disease-ghana/input/Corn/Corn__C... (4080, 1836)
1
2 /content/crop-disease-ghana/input/Corn/Corn__C... (4080, 1836)
1
3 /content/crop-disease-ghana/input/Corn/Corn__C... (4080, 1836)
1
4 /content/crop-disease-ghana/input/Corn/Corn__C... (4080, 1836)
1

```

```

# Crop type and disease distribution

```

```


plt.figure(figsize = (10,10))
sns.countplot(x = labels_df['crop'],hue = labels_df['disease'],palette
= 'Set2')

```

```

<Axes: xlabel='crop', ylabel='count'>

```





```
# Drop duplicates
no_duplicates = labels_df.drop_duplicates(subset =
'img_path',ignore_index=True)
no_duplicates.reset_index()
no_duplicates.shape

(6919, 14)

# Crop type and disease distribution
plt.figure(figsize = (10,10))
sns.countplot(x = no_duplicates['crop'],hue =
no_duplicates['disease'],palette = 'Set2')

<Axes: xlabel='crop', ylabel='count'>
```



```
# Disease value counts
no_duplicates['disease'].value_counts()
```

```
Corn Cercospora Leaf Spot    1072
Pepper Leaf Curl              763
Pepper Bacterial Spot        716
Tomato Leaf Blight           696
Tomato Septoria              579
Corn Streak                   467
Pepper Healthy               408
Corn Common Rust             348
Tomato Healthy               331
Pepper Leaf Blight           312
```

Pepper Fusarium	306
Corn Healthy	296
Pepper Cercospora	203
Pepper Septoria	169
Tomato Fusarium	126
Tomato Leaf Curl	68
Tomato Bacterial Spot	59

Name: disease, dtype: int64

Data Analysis

Graph shows class imbalance in selected dataset. There are categories such as Corn Cercospora Leaf, Pepper Leaf Curl, Pepper Bacteria Spot and others are contributing to majority of sample data. This will lead us in our next section where we addressed class imbalances.

Data Augmentation

In order to address class imbalance problem, there should be a combination of resampling and data augmentation.

```
# Modifed from:
# https://medium.com/@a.karazhay/guide-augment-images-and-multiple-bounding-boxes-for-deep-learning-in-4-steps-with-the-notebook-9b263e414dac
# Augments bounding box coordiantes along with their respective images

# Function to convert BoundingBoxesOnImage object into DataFrame
def bbs_obj_to_df(bbs_object):
    # Convert bounding box object to array
    bbs_array = bbs_object.to_xyxy_array()
    # Convert array into a DataFrame ['xmin', 'ymin', 'xmax', 'ymax']
    columns
    df_bbs = pd.DataFrame(bbs_array, columns=['xmin', 'ymin', 'xmax', 'ymax'])
    return df_bbs

def image_aug(df, aug_images_path, image_prefix, augmentor, ref_df):
    # Create data frame to store augmented image info
    aug_bbs_xy = pd.DataFrame(columns=labels_df.columns)
    grouped = df.groupby('filename')
    # Group data by filename
    for filename, i in zip(df['filename'].unique(), range(len(df))):
        group_df = grouped.get_group(filename)
        group_df = group_df.reset_index()
        group_df = group_df.drop(['index'], axis=1)
        # Read image
        image = imageio.imread(ref_df['img_path'][i])
        # Find bounding box coordinates put and put them into an array
```

```

        bb_array = group_df.drop(['filename', 'disease', 'crop',
                                'width', 'height', 'depth',
                                'ann_path',
                                'img_path', 'label', 'shape'], axis=1).values
        # Pass bounding box coordinates to Img Aug
        bbs = BoundingBoxesOnImage.from_xyxy_array(bb_array,
                                                    shape=image.shape)

        # Apply augmentation on image and on the bounding boxes
        image_aug, bbs_aug = augmentor(image=image,
                                         bounding_boxes=bbs)
        # Disregard boxes which have fallen out of image
        bbs_aug = bbs_aug.remove_out_of_image()
        # Clip boxes which are partially outside of image
        bbs_aug = bbs_aug.clip_out_of_image()

        # Don't perform any actions with the image if there are no
        bounding boxes left in it
        if re.findall('Image...', str(bbs_aug)) == ['Image([])']:
            pass
        else:
            # Write augmented image to a file
            os.chdir(aug_images_path)
            name = image_prefix+'_'+filename
            cv2.imwrite(name, image_aug)
            # Create a data frame with augmented values of image width
            and height
            info_df = group_df.drop(['xmin', 'ymin', 'xmax', 'ymax'],
                                     axis=1)
            for index, _ in info_df.iterrows():
                info_df.at[index, 'width'] = image_aug.shape[1]
                info_df.at[index, 'height'] = image_aug.shape[0]
            # Rename filenames by adding the predefined prefix
            info_df['filename'] = info_df['filename'].apply(lambda x:
image_prefix+'_'+x)
            # Create a data frame with augmented bounding boxes
            coordinates using the function we created earlier
            bbs_df = bbs_obj_to_df(bbs_aug)
            # Concat all new augmented info into new data frame
            aug_df = pd.concat([info_df, bbs_df], axis=1)
            # Append rows to aug_bbs_xy data frame
            aug_bbs_xy = pd.concat([aug_bbs_xy, aug_df])

        # Return dataframe with updated images and bounding boxes
        annotations
        aug_bbs_xy = aug_bbs_xy.reset_index()
        aug_bbs_xy = aug_bbs_xy.drop(['index'], axis=1)
        return aug_bbs_xy

```

```

# Augmentors to use
aug_flip = iaa.Sequential([
    iaa.Fliplr(1)
])
aug_rotate = iaa.Sequential([
    iaa.Affine(rotate=(-60, 60))
])

aug_blur = iaa.Sequential([
    iaa.GaussianBlur(sigma=(2.0, 3.0))
])

aug_noise = iaa.Sequential([
    iaa.AdditiveGaussianNoise(scale=(0.03*255, 0.05*255))
])

# Going to augment rare classes the append augmented data to labels_df
tlc_df = labels_df[labels_df['disease']=='Tomato Leaf Curl'].reset_index()
tbs_df = labels_df[labels_df['disease']=='Tomato Bacterial Spot'].reset_index()
pc_df = labels_df[labels_df['disease']=='Pepper Cercospora'].reset_index()
ps_df = labels_df[labels_df['disease']=='Pepper Septoria'].reset_index()
tf_df = labels_df[labels_df['disease']=='Tomato Fusarium'].reset_index()

# Make folders to store augmented images
!mkdir Tomato_Leaf_Curl_Aug
!mkdir Tomato_Bacterial_Spot_Aug
!mkdir Pepper_Cercospora_Aug
!mkdir Pepper_Septoria_Aug
!mkdir Tomato_Fusarium_Aug

# Create augmented data set
def make_aug(df,aug_dir):
    aug_list = [aug_flip,aug_rotate,aug_blur,aug_noise]

    aug_df = image_aug(df[labels_df.columns],
                        aug_dir, 'flip', aug_list[0],df)
    label = ['rotate','blur','noise']
    for a,n in zip(aug_list[1:],label):
        aug_df = aug_df.append(image_aug(df[labels_df.columns],
                                          aug_dir, n, a,df))

    # Drop rows with null
    aug_df = aug_df.dropna()
    return aug_df

```

```
# Make augmented datasets
```

```
tlc_aug =
make_aug(tlc_df, '/content/Tomato_Leaf_Curl_Aug').reset_index()
tbs_aug =
make_aug(tbs_df, '/content/Tomato_Bacterial_Spot_Aug').reset_index()
pc_aug =
make_aug(pc_df, '/content/Pepper_Cercospora_Aug').reset_index()
ps_aug = make_aug(ps_df, '/content/Pepper_Septoria_Aug').reset_index()
tf_aug = make_aug(tf_df, '/content/Tomato_Fusarium_Aug').reset_index()

tlc_aug.head()
```

	index	filename	disease	crop	width	height
depth \						
0	0	flip_IMG_3011.jpeg	Tomato Leaf Curl	Tomato	4032	3024
3						
1	1	flip_IMG_3011.jpeg	Tomato Leaf Curl	Tomato	4032	3024
3						
2	2	flip_IMG_3011.jpeg	Tomato Leaf Curl	Tomato	4032	3024
3						
3	3	flip_IMG_3011.jpeg	Tomato Leaf Curl	Tomato	4032	3024
3						
4	4	flip_IMG_3011.jpeg	Tomato Leaf Curl	Tomato	4032	3024
3						

	xmin	ymin	xmax	ymax	\
0	1533.429932	1586.780884	1840.773926	2167.487305	
1	618.202393	1597.042847	1096.607178	2033.170288	
2	2577.438477	902.496277	2912.614014	1354.034424	
3	1315.911865	1372.759277	1736.795410	1517.977295	
4	2185.676270	1207.585449	2512.000732	1416.671509	

	ann_path	\
0	/content/crop-disease-ghana/input/Tomato/Tomat...	
1	/content/crop-disease-ghana/input/Tomato/Tomat...	
2	/content/crop-disease-ghana/input/Tomato/Tomat...	
3	/content/crop-disease-ghana/input/Tomato/Tomat...	
4	/content/crop-disease-ghana/input/Tomato/Tomat...	

	img_path	shape
label		
0	/content/crop-disease-ghana/input/Tomato/Tomat...	(4032, 3024)
16		
1	/content/crop-disease-ghana/input/Tomato/Tomat...	(4032, 3024)
16		
2	/content/crop-disease-ghana/input/Tomato/Tomat...	(4032, 3024)
16		
3	/content/crop-disease-ghana/input/Tomato/Tomat...	(4032, 3024)
16		

```

4 /content/crop-disease-ghana/input/Tomato/Tomat... (4032, 3024)
16

# Display augmented image
p = '/content/Tomato_Leaf_Curl_Aug/' + tlc_aug['filename'][0]
test_img=cv2.imread(p)
def display_image(test_img):
    h, w = test_img.shape[0:2]
    neww = 300
    newh = int(neww*(h/w))
    test_img = cv2.resize(test_img, (neww, newh))
    cv2.imshow(test_img)
    cv2.waitKey(0)

display_image(test_img)

```



```

# Replace image paths
tlc_aug['img_path'] = '/content/Tomato_Leaf_Curl_Aug/' +
tlc_aug['filename']
tbs_aug['img_path'] = '/content/Tomato_Bacterial_Spot_Aug/' +
tbs_aug['filename']
pc_aug['img_path'] = '/content/Pepper_Cercospora_Aug/' +
pc_aug['filename']
ps_aug['img_path'] = '/content/Pepper_Septoria_Aug/' +
ps_aug['filename']
tf_aug['img_path'] = '/content/Tomato_Fusarium_Aug/' +
tf_aug['filename']

# Drop index column
tlc_aug = tlc_aug.drop(labels = 'index',axis = 1)
tbs_aug = tbs_aug.drop(labels = 'index',axis = 1)

```



```

pc_aug = pc_aug.drop(labels = 'index',axis = 1)
ps_aug = ps_aug.drop(labels = 'index',axis = 1)
tf_aug = tf_aug.drop(labels = 'index',axis = 1)

# Append to labels_df
labels_df = labels_df.append(tlc_aug)
labels_df = labels_df.append(tbs_aug)
labels_df = labels_df.append(pc_aug)
labels_df = labels_df.append(ps_aug)
labels_df = labels_df.append(tf_aug)
labels_df = labels_df.reset_index()

labels_df.shape

(70259, 15)

# Drop duplicates
no_duplicates = labels_df.drop_duplicates(subset =
'filename',ignore_index=True)
no_duplicates.shape

(9191, 15)

no_duplicates['disease'].value_counts()

Pepper Cercospora      1096
Corn Cercospora Leaf Spot  1021
Pepper Septoria        835
Pepper Leaf Curl       760
Pepper Bacterial Spot   706
Tomato Leaf Blight      632
Tomato Fusarium         570
Corn Streak            467
Tomato Leaf Curl       447
Tomato Bacterial Spot   412
Tomato Septoria         409
Pepper Healthy         366
Tomato Healthy          331
Pepper Leaf Blight     312
Pepper Fusarium        306
Corn Healthy           280
Corn Common Rust       241
Name: disease, dtype: int64

# Create final balanced dataset
balanced = pd.DataFrame(columns = labels_df.columns)
for i in no_duplicates['disease'].unique():
    df = no_duplicates[no_duplicates['disease'] == i]
    if len(df)>200:
        samp = df.sample(n = 200,random_state = 24)
        balanced = balanced.append(samp)

```

```

else:
    balanced = balanced.append(df)

# Using less data because of computational constraints
balanced['disease'].value_counts()

Corn Cercospora Leaf Spot      200
Pepper Fusarium                 200
Tomato Leaf Curl                200
Tomato Fusarium                 200
Tomato Septoria                 200
Tomato Bacterial Spot           200
Tomato Leaf Blight              200
Pepper Septoria                 200
Pepper Healthy                  200
Corn Common Rust                200
Pepper Bacterial Spot           200
Pepper Leaf Blight              200
Pepper Cercospora               200
Pepper Leaf Curl                200
Corn Streak                     200
Corn Healthy                    200
Tomato Healthy                   200
Name: disease, dtype: int64

# Create final balanced df
balanced_df = pd.DataFrame(columns = balanced.columns)
for i in np.unique(balanced['filename']):
    df = labels_df[labels_df['filename'] == i]
    balanced_df = balanced_df.append(df,ignore_index = True)
#balanced_df = balanced_df.reset_index()
#balanced_df = balanced_df.drop(labels = ['index','level_0'],axis = 1)
balanced_df.head()

```

	index	filename	disease	crop	width
height \					
0	21667	2010-01-01%2000.06.21.jpg	Pepper Healthy	Pepper	1920
2560					
1	21668	2010-01-01%2000.06.21.jpg	Pepper Healthy	Pepper	1920
2560					
2	19004	2010-01-01%2000.11.06.jpg	Pepper Cercospora	Pepper	1920
2560					
3	19005	2010-01-01%2000.11.06.jpg	Pepper Cercospora	Pepper	1920
2560					
4	19006	2010-01-01%2000.11.06.jpg	Pepper Cercospora	Pepper	1920
2560					
depth	xmin	ymin	xmax	ymax	\
0	3	938.970994	1550.683230	1672.769337	2274.989648

1	3	281.933702	1781.283644	611.042818	2365.175983
2	3	1199.818892	1273.706070	1272.400568	1334.632588
3	3	952.734375	1139.371672	1015.887784	1173.418530
4	3	1037.791193	993.439830	1091.313920	1037.018104

	ann_path	\
0	/content/crop-disease-ghana/input/Pepper/Peppe...	
1	/content/crop-disease-ghana/input/Pepper/Peppe...	
2	/content/crop-disease-ghana/input/Pepper/Peppe...	
3	/content/crop-disease-ghana/input/Pepper/Peppe...	
4	/content/crop-disease-ghana/input/Pepper/Peppe...	

	img_path	shape
label		
0	/content/crop-disease-ghana/input/Pepper/Peppe...	(1920, 2560)
8		
1	/content/crop-disease-ghana/input/Pepper/Peppe...	(1920, 2560)
8		
2	/content/crop-disease-ghana/input/Pepper/Peppe...	(1920, 2560)
6		
3	/content/crop-disease-ghana/input/Pepper/Peppe...	(1920, 2560)
6		
4	/content/crop-disease-ghana/input/Pepper/Peppe...	(1920, 2560)
6		

Dropping unnecessary columns

```
balanced_df = balanced_df.drop(labels=['index'],axis = 1)
balanced_df.head()
```

	filename	disease	crop	width	height
depth	\				
0	2010-01-01%2000.06.21.jpg	Pepper Healthy	Pepper	1920	2560
3					
1	2010-01-01%2000.06.21.jpg	Pepper Healthy	Pepper	1920	2560
3					
2	2010-01-01%2000.11.06.jpg	Pepper Cercospora	Pepper	1920	2560
3					
3	2010-01-01%2000.11.06.jpg	Pepper Cercospora	Pepper	1920	2560
3					
4	2010-01-01%2000.11.06.jpg	Pepper Cercospora	Pepper	1920	2560
3					

	xmin	ymin	xmax	ymax	\
0	938.970994	1550.683230	1672.769337	2274.989648	
1	281.933702	1781.283644	611.042818	2365.175983	
2	1199.818892	1273.706070	1272.400568	1334.632588	
3	952.734375	1139.371672	1015.887784	1173.418530	
4	1037.791193	993.439830	1091.313920	1037.018104	

	ann_path	\
--	----------	---

```

0 /content/crop-disease-ghana/input/Pepper/Peppe...
1 /content/crop-disease-ghana/input/Pepper/Peppe...
2 /content/crop-disease-ghana/input/Pepper/Peppe...
3 /content/crop-disease-ghana/input/Pepper/Peppe...
4 /content/crop-disease-ghana/input/Pepper/Peppe...

```

	img_path	shape
label		
0	/content/crop-disease-ghana/input/Pepper/Peppe...	(1920, 2560)
8		
1	/content/crop-disease-ghana/input/Pepper/Peppe...	(1920, 2560)
8		
2	/content/crop-disease-ghana/input/Pepper/Peppe...	(1920, 2560)
6		
3	/content/crop-disease-ghana/input/Pepper/Peppe...	(1920, 2560)
6		
4	/content/crop-disease-ghana/input/Pepper/Peppe...	(1920, 2560)
6		

```

df_balanced = pd.DataFrame(columns = labels_df.columns)
for i in np.unique(balanced_df['disease']):
    df = balanced_df[balanced_df['disease'] == i]
    if len(df)>200:
        samp = df.sample(n = 200,random_state = 24)
        df_balanced = df_balanced.append(samp,ignore_index = True)
    else:
        df_balanced = df_balanced.append(df,ignore_index = True)

df_balanced = df_balanced.drop(labels = ['index'],axis = 1)
df_balanced.head()

```

	filename	disease	crop	width
height depth \				
0	20230526_114545.jpg	Corn Cercospora Leaf Spot	Corn	4080
1836	3			
1	20230526_103506(0).jpg	Corn Cercospora Leaf Spot	Corn	4080
1836	3			
2	20230525_111356(0).jpg	Corn Cercospora Leaf Spot	Corn	4080
1836	3			
3	20230524_110802.jpg	Corn Cercospora Leaf Spot	Corn	4080
1836	3			
4	20230526_102705.jpg	Corn Cercospora Leaf Spot	Corn	4080
1836	3			

	xmin	ymin	xmax	ymax	\
0	2089.895178	897.998932	2172.578616	1003.483405	
1	2717.148847	592.949243	2822.641509	672.775330	
2	2941.199173	1275.034169	3076.168160	1407.181183	
3	2857.733711	753.781004	2941.529745	826.064469	
4	2976.603774	1257.216324	3062.138365	1348.446137	

```

                                ann_path  \
0  /content/crop-disease-ghana/input/Corn/Corn__C...
1  /content/crop-disease-ghana/input/Corn/Corn__C...
2  /content/crop-disease-ghana/input/Corn/Corn__C...
3  /content/crop-disease-ghana/input/Corn/Corn__C...
4  /content/crop-disease-ghana/input/Corn/Corn__C...

                                img_path      shape
label
0  /content/crop-disease-ghana/input/Corn/Corn__C...  (4080, 1836)
1
1  /content/crop-disease-ghana/input/Corn/Corn__C...  (4080, 1836)
1
2  /content/crop-disease-ghana/input/Corn/Corn__C...  (4080, 1836)
1
3  /content/crop-disease-ghana/input/Corn/Corn__C...  (4080, 1836)
1
4  /content/crop-disease-ghana/input/Corn/Corn__C...  (4080, 1836)
1

```

Multi-class object detection and bounding box regression

A single forward pass of our multi-class object detector will result in:

1. The predicted bounding box coordinates of the object in the image
2. The predicted class label of the object in the image

Model Design

Branch #1: A regression layer set, just like in the single-class object detection case

Branch #2: An additional layer set, this one with a softmax classifier used to predict class labels

```

# Initialize parameters

width=224
height=224
depth= 3
epoch_= 25
BS = 32
default_image_size = tuple((224, 224))
image_size = 0
INIT_LR = 1e-3

```

We use `resize` and `img_to_array` method of open cv to ensure that our image size is 224x 224 pixels for training with VGG16 followed by converting to array format

```
def convert_image_to_array(image_dir):  
    try:  
        image = cv2.imread(image_dir)  
        if image is not None :  
            image = cv2.resize(image, default_image_size)  
            return img_to_array(image)  
        else :  
            return np.array([])  
    except Exception as e:  
        print(f"Error : {e}")  
        return None  
  
df_balanced.shape  
  
(3400, 14)  
  
# scale the bounding box coordinates relative to the spatial  
dimensions of the input image  
new_xmin = []  
new_ymin = []  
new_xmax = []  
new_ymax = []  
  
for i in range(df_balanced.shape[0]):  
    w=df_balanced['width'][i]  
    h=df_balanced['height'][i]  
    xmin=df_balanced['xmin'][i]  
    ymin=df_balanced['ymin'][i]  
    xmax=df_balanced['xmax'][i]  
    ymax=df_balanced['ymax'][i]  
    new_xmin.append(float(xmin) / w)  
    new_ymin.append(float(ymin) / h)  
    new_xmax.append(float(xmax) / w)  
    new_ymax.append(float(ymax) / h)  
  
df_balanced['xmin'] = new_xmin  
df_balanced['ymin'] = new_ymin  
df_balanced['xmax'] = new_xmax  
df_balanced['ymax'] = new_ymax  
df_balanced.head()
```

	filename	disease	crop	width
height depth \				
0	20230526_114545.jpg	Corn Cercospora Leaf Spot	Corn	4080
1836	3			
1	20230526_103506(0).jpg	Corn Cercospora Leaf Spot	Corn	4080
1836	3			

```

2 20230525_111356(0).jpg  Corn Cercospora Leaf Spot  Corn  4080
1836      3
3 20230524_110802.jpg  Corn Cercospora Leaf Spot  Corn  4080
1836      3
4 20230526_102705.jpg  Corn Cercospora Leaf Spot  Corn  4080
1836      3

```

```

      xmin      ymin      xmax      ymax  \
0  0.512229  0.489106  0.532495  0.546560
1  0.665968  0.322957  0.691824  0.366435
2  0.720882  0.694463  0.753963  0.766439
3  0.700425  0.410556  0.720963  0.449926
4  0.729560  0.684758  0.750524  0.734448

```

```

                                ann_path  \
0  /content/crop-disease-ghana/input/Corn/Corn__C...
1  /content/crop-disease-ghana/input/Corn/Corn__C...
2  /content/crop-disease-ghana/input/Corn/Corn__C...
3  /content/crop-disease-ghana/input/Corn/Corn__C...
4  /content/crop-disease-ghana/input/Corn/Corn__C...

```

```

                                img_path      shape
label
0  /content/crop-disease-ghana/input/Corn/Corn__C...  (4080, 1836)
1
1  /content/crop-disease-ghana/input/Corn/Corn__C...  (4080, 1836)
1
2  /content/crop-disease-ghana/input/Corn/Corn__C...  (4080, 1836)
1
3  /content/crop-disease-ghana/input/Corn/Corn__C...  (4080, 1836)
1
4  /content/crop-disease-ghana/input/Corn/Corn__C...  (4080, 1836)
1

```

Finally, we populate those four lists that we initialized : (1) image_list (array format), (2) ImagePaths_list (3) label_list(Disease catagories), and (4) bboxes_list(xmin ymin, xmax, ymax co-ordinates)

```

image_list, label_list, bboxes_list, imagePaths_list = [], [], [], []

indx=0
i=0
try:
    for img_dir in df_balanced['img_path']:
        #for i in range(2):
            #img_dir=
            f"{root_dir}/{plant_folder}/{plant_disease_folder}/{image}"
            if img_dir.endswith(".jpg") == True or img_dir.endswith(".JPG") ==
True:
                image_list.append(convert_image_to_array(img_dir))

```



```

        imagePath_list.append(df_balanced['img_path'][indx])
        label_list.append(df_balanced['disease'][indx])
        bboxes_list.append((df_balanced['xmin'][indx],
df_balanced['ymin'][indx], df_balanced['xmax'][indx],
df_balanced['ymax'][indx]))
        indx=indx+1
except Exception as e:
    print(f"Error : {e}")

data, labels, bboxes, imagePath_list = [], [], [], []
data = np.array(image_list, dtype="float32") / 255.0
labels = np.array(label_list)
bboxes = np.array(bboxes_list, dtype="float32")
imagePaths = np.array(imagePaths_list)

print (data[0], labels[0], bboxes[0], imagePath_list[0])

[[[0.23921569 0.73333335 0.50980395]
 [0.35686275 0.8156863 0.60784316]
 [0.33333334 0.79607844 0.59607846]
 ...
 [0. 0.36862746 0.23137255]
 [0. 0.3137255 0.14509805]
 [0. 0.2784314 0.10980392]]

 [[0.1254902 0.5921569 0.36078432]
 [0.12941177 0.5882353 0.35686275]
 [0.25490198 0.75686276 0.5411765 ]
 ...
 [0. 0.30980393 0.11764706]
 [0. 0.27450982 0.10588235]
 [0. 0.26666668 0.09411765]]

 [[0.08627451 0.59607846 0.34901962]
 [0.12156863 0.59607846 0.34901962]
 [0.13725491 0.6 0.36862746]
 ...
 [0. 0.27450982 0.10588235]
 [0. 0.26666668 0.10588235]
 [0. 0.24313726 0.08235294]]

 ...

 [[0.5764706 0.73333335 0.60784316]
 [0.49411765 0.69411767 0.54509807]
 [0.39607844 0.62352943 0.4627451 ]
 ...
 [0.1254902 0.22745098 0.25490198]
 [0.08235294 0.19607843 0.21568628]
 [0.07843138 0.19215687 0.21176471]]

```

```

[[0.49019608 0.69803923 0.5529412 ]
 [0.36078432 0.5921569  0.42745098]
 [0.37254903 0.6039216  0.43137255]
 ...
 [0.09019608 0.2         0.22745098]
 [0.09411765 0.20784314 0.22745098]
 [0.10196079 0.20392157 0.22745098]]

[[0.3137255  0.56078434 0.39607844]
 [0.3529412  0.5882353  0.4117647 ]
 [0.33333334 0.5568628  0.38039216]
 ...
 [0.08235294 0.19607843 0.22352941]
 [0.07843138 0.1882353  0.21568628]
 [0.14901961 0.26666668 0.2901961 ]]] Corn Cercospora Leaf Spot
[0.5122292  0.48910618 0.5324948  0.5465596 ] /content/crop-disease-
ghana/input/Corn/Corn__Cercospora_Leaf_Spot/images/20230526_114545.jpg

```

One-hot encoded labels

```

label_binarizer = LabelBinarizer()
if labels is not None:
    binarylabels = label_binarizer.fit_transform(labels)
    pickle.dump(label_binarizer, open('label_transform.pkl', 'wb'))
    n_classes = len(label_binarizer.classes_)

print(labels.shape)

(2479,)

split = train_test_split(data, binarylabels, bboxes, test_size=0.20,
random_state=42)
# unpack the data split
(trainImages, testImages) = split[:2]
(trainLabels, testLabels) = split[2:4]
(trainBBoxes, testBBoxes) = split[4:6]
#(trainPaths, testPaths) = split[6:]

import os
import tensorflow as tf
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'

```

Pre-Trained Model VGG16

VGG16: The CNN architecture to serve as the base network which we'll (1) modify for multi-class bounding box regression and (2) then fine-tune on our dataset

VGG16 network is loaded with weights pre-trained on the ImageNet dataset. We leave off the fully-connected layer head (include_top=False), since we will be constructing a new layer head responsible for multi-output prediction (i.e., class label and bounding box location).

```

# load the VGG16 network, ensuring the head FC layers are left off
vgg = VGG16(weights="imagenet", include_top=False,
input_tensor=Input(shape=(224, 224, 3)))
# freeze all VGG layers so they will *not* be updated during the
# training process
vgg.trainable = False
# flatten the max-pooling output of VGG
flatten = vgg.output
flatten = Flatten()(flatten)

```

Flatten the output of the network so we can construct our new layer

(1) The first branch, bboxHead, is responsible for predicting the bounding box (x, y)-coordinates of the object in the image.

(2) Our second branch, softmaxHead, is responsible for predicting the class label of the detected object

```

# construct a fully-connected layer header to output the predicted
# bounding box coordinates
bboxHead = Dense(128, activation="relu")(flatten)
bboxHead = Dense(64, activation="relu")(bboxHead)
bboxHead = Dense(32, activation="relu")(bboxHead)
bboxHead = Dense(4, activation="sigmoid", name="bounding_box")
(bboxHead)
# construct a second fully-connected layer head, this one to predict
# the class label
softmaxHead = Dense(512, activation="relu")(flatten)
softmaxHead = Dropout(0.5)(softmaxHead)
softmaxHead = Dense(512, activation="relu")(softmaxHead)
softmaxHead = Dropout(0.5)(softmaxHead)
softmaxHead = Dense(len(label_binarizer.classes_),
activation="softmax", name="class_label")(softmaxHead)
# put together our model which accept an input image and then output
# bounding box coordinates and a class label
model = Model(
    inputs=vgg.input,
    outputs=(bboxHead, softmaxHead))

```

We are using categorical cross-entropy for our class label branch and mean squared error for our bounding box regression head.

```

# define a dictionary to set the loss methods -- categorical
# cross-entropy for the class label head and mean absolute error
# for the bounding box head
losses = {
    "bounding_box": "mean_squared_error",
    "class_label": "categorical_crossentropy"
}

```

```

    #"bounding_box": "mean_squared_error",
}
# define a dictionary that specifies the weights per loss (both the
# class label and bounding box outputs will receive equal weight)
lossWeights = {
    "bounding_box": 1.0,
    "class_label": 1.0
    #"bounding_box": 1.0
}
# initialize the optimizer, compile the model, and show the model
# summary
opt = Adam(learning_rate=INIT_LR)
modell.compile(loss=losses, optimizer=opt, metrics=["accuracy"],
loss_weights=lossWeights)
print(modell.summary())

```

Model: "model"

Layer (type)	Output Shape	Param #
Connected to		
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
=====		
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
['input_1[0][0]']		
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
['block1_conv1[0][0]']		
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
['block1_conv2[0][0]']		
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
['block1_pool[0][0]']		
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
['block2_conv1[0][0]']		
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
['block2_conv2[0][0]']		

block3_conv1 (Conv2D) ['block2_pool[0][0]']	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D) ['block3_conv1[0][0]']	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D) ['block3_conv2[0][0]']	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D) ['block3_conv3[0][0]']	(None, 28, 28, 256)	0
block4_conv1 (Conv2D) ['block3_pool[0][0]']	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D) ['block4_conv1[0][0]']	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D) ['block4_conv2[0][0]']	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D) ['block4_conv3[0][0]']	(None, 14, 14, 512)	0
block5_conv1 (Conv2D) ['block4_pool[0][0]']	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D) ['block5_conv1[0][0]']	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D) ['block5_conv2[0][0]']	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D) ['block5_conv3[0][0]']	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0

```
['block5_pool[0][0]']
```

dense_3 (Dense)	(None, 512)	1284556
['flatten[0][0]']		8

dense (Dense)	(None, 128)	3211392
['flatten[0][0]']		

dropout (Dropout)	(None, 512)	0
['dense_3[0][0]']		

dense_1 (Dense)	(None, 64)	8256
['dense[0][0]']		

dense_4 (Dense)	(None, 512)	262656
['dropout[0][0]']		

dense_2 (Dense)	(None, 32)	2080
['dense_1[0][0]']		

dropout_1 (Dropout)	(None, 512)	0
['dense_4[0][0]']		

bounding_box (Dense)	(None, 4)	132
['dense_2[0][0]']		

class_label (Dense)	(None, 13)	6669
['dropout_1[0][0]']		

```
=====
```

```
Total params: 31051441 (118.45 MB)
Trainable params: 16336753 (62.32 MB)
Non-trainable params: 14714688 (56.13 MB)
```

None

```

# construct a dictionary for our target training outputs
trainTargets = {
    "bounding_box": trainBBoxes,
    "class_label": trainLabels
    #"bounding_box": trainBBoxes
}
# construct a second dictionary, this one for our target testing
# outputs
testTargets = {
    "bounding_box": testBBoxes,
    "class_label": testLabels
    #"bounding_box": testBBoxes
}

```

Training multi-class object detector

Our architecture has two branches in the layer head — the first branch to predict the bounding box coordinates and the second to predict the class label of the detected object

```

# train the network for bounding box regression and class label
# prediction
print("[INFO] training model...")
H = model.fit(
    trainImages, trainTargets,
    validation_data=(testImages, testTargets),
    batch_size=BS,
    epochs=50,
    #class_weight=class_weights,
    verbose=1)

```

[INFO] training model...

Epoch 1/50

62/62 [=====] - 19s 139ms/step - loss: 3.0276 - bounding_box_loss: 0.0585 - class_label_loss: 2.9690 - bounding_box_accuracy: 0.4347 - class_label_accuracy: 0.1906 - val_loss: 1.9419 - val_bounding_box_loss: 0.0535 - val_class_label_loss: 1.8884 - val_bounding_box_accuracy: 0.5262 - val_class_label_accuracy: 0.3407

Epoch 2/50

62/62 [=====] - 4s 67ms/step - loss: 2.1131 - bounding_box_loss: 0.0487 - class_label_loss: 2.0644 - bounding_box_accuracy: 0.4851 - class_label_accuracy: 0.2481 - val_loss: 1.9010 - val_bounding_box_loss: 0.0539 - val_class_label_loss: 1.8471 - val_bounding_box_accuracy: 0.5302 - val_class_label_accuracy: 0.3770

Epoch 3/50

62/62 [=====] - 4s 65ms/step - loss: 1.9877 - bounding_box_loss: 0.0459 - class_label_loss: 1.9418 -

bounding_box_accuracy: 0.4634 - class_label_accuracy: 0.2960 -
val_loss: 1.6535 - val_bounding_box_loss: 0.0586 -
val_class_label_loss: 1.5949 - val_bounding_box_accuracy: 0.4617 -
val_class_label_accuracy: 0.4516
Epoch 4/50
62/62 [=====] - 4s 69ms/step - loss: 1.8881 -
bounding_box_loss: 0.0455 - class_label_loss: 1.8425 -
bounding_box_accuracy: 0.4912 - class_label_accuracy: 0.3167 -
val_loss: 1.5108 - val_bounding_box_loss: 0.0531 -
val_class_label_loss: 1.4577 - val_bounding_box_accuracy: 0.5323 -
val_class_label_accuracy: 0.5161
Epoch 5/50
62/62 [=====] - 4s 68ms/step - loss: 1.7252 -
bounding_box_loss: 0.0421 - class_label_loss: 1.6831 -
bounding_box_accuracy: 0.5214 - class_label_accuracy: 0.3732 -
val_loss: 1.5221 - val_bounding_box_loss: 0.0561 -
val_class_label_loss: 1.4659 - val_bounding_box_accuracy: 0.3569 -
val_class_label_accuracy: 0.5181
Epoch 6/50
62/62 [=====] - 5s 76ms/step - loss: 1.7345 -
bounding_box_loss: 0.0397 - class_label_loss: 1.6948 -
bounding_box_accuracy: 0.5265 - class_label_accuracy: 0.3792 -
val_loss: 1.3897 - val_bounding_box_loss: 0.0550 -
val_class_label_loss: 1.3348 - val_bounding_box_accuracy: 0.4919 -
val_class_label_accuracy: 0.5181
Epoch 7/50
62/62 [=====] - 5s 81ms/step - loss: 1.6394 -
bounding_box_loss: 0.0369 - class_label_loss: 1.6025 -
bounding_box_accuracy: 0.6031 - class_label_accuracy: 0.4014 -
val_loss: 1.3815 - val_bounding_box_loss: 0.0548 -
val_class_label_loss: 1.3267 - val_bounding_box_accuracy: 0.4133 -
val_class_label_accuracy: 0.5665
Epoch 8/50
62/62 [=====] - 4s 70ms/step - loss: 1.5505 -
bounding_box_loss: 0.0354 - class_label_loss: 1.5151 -
bounding_box_accuracy: 0.6051 - class_label_accuracy: 0.4216 -
val_loss: 1.3537 - val_bounding_box_loss: 0.0609 -
val_class_label_loss: 1.2928 - val_bounding_box_accuracy: 0.4093 -
val_class_label_accuracy: 0.5403
Epoch 9/50
62/62 [=====] - 4s 66ms/step - loss: 1.5244 -
bounding_box_loss: 0.0350 - class_label_loss: 1.4895 -
bounding_box_accuracy: 0.6319 - class_label_accuracy: 0.4231 -
val_loss: 1.3108 - val_bounding_box_loss: 0.0566 -
val_class_label_loss: 1.2542 - val_bounding_box_accuracy: 0.5282 -
val_class_label_accuracy: 0.5685
Epoch 10/50
62/62 [=====] - 4s 67ms/step - loss: 1.5147 -
bounding_box_loss: 0.0329 - class_label_loss: 1.4817 -

bounding_box_accuracy: 0.6157 - class_label_accuracy: 0.4276 -
val_loss: 1.4049 - val_bounding_box_loss: 0.0595 -
val_class_label_loss: 1.3454 - val_bounding_box_accuracy: 0.4315 -
val_class_label_accuracy: 0.5323
Epoch 11/50
62/62 [=====] - 5s 79ms/step - loss: 1.4535 -
bounding_box_loss: 0.0318 - class_label_loss: 1.4217 -
bounding_box_accuracy: 0.6490 - class_label_accuracy: 0.4629 -
val_loss: 1.2777 - val_bounding_box_loss: 0.0586 -
val_class_label_loss: 1.2190 - val_bounding_box_accuracy: 0.4980 -
val_class_label_accuracy: 0.5867
Epoch 12/50
62/62 [=====] - 5s 75ms/step - loss: 1.3957 -
bounding_box_loss: 0.0303 - class_label_loss: 1.3654 -
bounding_box_accuracy: 0.6475 - class_label_accuracy: 0.4584 -
val_loss: 1.2728 - val_bounding_box_loss: 0.0625 -
val_class_label_loss: 1.2103 - val_bounding_box_accuracy: 0.5181 -
val_class_label_accuracy: 0.6149
Epoch 13/50
62/62 [=====] - 4s 66ms/step - loss: 1.4315 -
bounding_box_loss: 0.0304 - class_label_loss: 1.4011 -
bounding_box_accuracy: 0.6727 - class_label_accuracy: 0.4493 -
val_loss: 1.2688 - val_bounding_box_loss: 0.0632 -
val_class_label_loss: 1.2056 - val_bounding_box_accuracy: 0.5101 -
val_class_label_accuracy: 0.6089
Epoch 14/50
62/62 [=====] - 5s 80ms/step - loss: 1.3934 -
bounding_box_loss: 0.0307 - class_label_loss: 1.3627 -
bounding_box_accuracy: 0.6803 - class_label_accuracy: 0.4695 -
val_loss: 1.2336 - val_bounding_box_loss: 0.0582 -
val_class_label_loss: 1.1753 - val_bounding_box_accuracy: 0.5242 -
val_class_label_accuracy: 0.6069
Epoch 15/50
62/62 [=====] - 4s 66ms/step - loss: 1.3835 -
bounding_box_loss: 0.0288 - class_label_loss: 1.3548 -
bounding_box_accuracy: 0.6899 - class_label_accuracy: 0.4755 -
val_loss: 1.3159 - val_bounding_box_loss: 0.0601 -
val_class_label_loss: 1.2559 - val_bounding_box_accuracy: 0.5222 -
val_class_label_accuracy: 0.6230
Epoch 16/50
62/62 [=====] - 5s 76ms/step - loss: 1.4599 -
bounding_box_loss: 0.0301 - class_label_loss: 1.4298 -
bounding_box_accuracy: 0.6662 - class_label_accuracy: 0.4690 -
val_loss: 1.3568 - val_bounding_box_loss: 0.0620 -
val_class_label_loss: 1.2948 - val_bounding_box_accuracy: 0.4294 -
val_class_label_accuracy: 0.5504
Epoch 17/50
62/62 [=====] - 4s 70ms/step - loss: 1.3438 -
bounding_box_loss: 0.0278 - class_label_loss: 1.3161 -

bounding_box_accuracy: 0.6798 - class_label_accuracy: 0.4771 -
val_loss: 1.2291 - val_bounding_box_loss: 0.0632 -
val_class_label_loss: 1.1659 - val_bounding_box_accuracy: 0.5081 -
val_class_label_accuracy: 0.6411
Epoch 18/50
62/62 [=====] - 4s 67ms/step - loss: 1.3262 -
bounding_box_loss: 0.0265 - class_label_loss: 1.2997 -
bounding_box_accuracy: 0.6949 - class_label_accuracy: 0.4927 -
val_loss: 1.2151 - val_bounding_box_loss: 0.0640 -
val_class_label_loss: 1.1511 - val_bounding_box_accuracy: 0.5202 -
val_class_label_accuracy: 0.6452
Epoch 19/50
62/62 [=====] - 5s 76ms/step - loss: 1.3230 -
bounding_box_loss: 0.0269 - class_label_loss: 1.2961 -
bounding_box_accuracy: 0.7196 - class_label_accuracy: 0.4987 -
val_loss: 1.1984 - val_bounding_box_loss: 0.0620 -
val_class_label_loss: 1.1364 - val_bounding_box_accuracy: 0.4980 -
val_class_label_accuracy: 0.5847
Epoch 20/50
62/62 [=====] - 4s 70ms/step - loss: 1.3395 -
bounding_box_loss: 0.0263 - class_label_loss: 1.3132 -
bounding_box_accuracy: 0.7237 - class_label_accuracy: 0.4791 -
val_loss: 1.2191 - val_bounding_box_loss: 0.0620 -
val_class_label_loss: 1.1572 - val_bounding_box_accuracy: 0.5222 -
val_class_label_accuracy: 0.6169
Epoch 21/50
62/62 [=====] - 5s 78ms/step - loss: 1.4042 -
bounding_box_loss: 0.0262 - class_label_loss: 1.3781 -
bounding_box_accuracy: 0.7126 - class_label_accuracy: 0.4705 -
val_loss: 1.2598 - val_bounding_box_loss: 0.0623 -
val_class_label_loss: 1.1975 - val_bounding_box_accuracy: 0.4899 -
val_class_label_accuracy: 0.6290
Epoch 22/50
62/62 [=====] - 5s 75ms/step - loss: 1.2860 -
bounding_box_loss: 0.0257 - class_label_loss: 1.2604 -
bounding_box_accuracy: 0.7090 - class_label_accuracy: 0.4801 -
val_loss: 1.1940 - val_bounding_box_loss: 0.0641 -
val_class_label_loss: 1.1299 - val_bounding_box_accuracy: 0.5181 -
val_class_label_accuracy: 0.6431
Epoch 23/50
62/62 [=====] - 5s 78ms/step - loss: 1.3103 -
bounding_box_loss: 0.0258 - class_label_loss: 1.2845 -
bounding_box_accuracy: 0.7171 - class_label_accuracy: 0.4786 -
val_loss: 1.1831 - val_bounding_box_loss: 0.0635 -
val_class_label_loss: 1.1195 - val_bounding_box_accuracy: 0.4657 -
val_class_label_accuracy: 0.6593
Epoch 24/50
62/62 [=====] - 5s 75ms/step - loss: 1.3063 -
bounding_box_loss: 0.0259 - class_label_loss: 1.2804 -

bounding_box_accuracy: 0.7161 - class_label_accuracy: 0.4922 -
val_loss: 1.2153 - val_bounding_box_loss: 0.0638 -
val_class_label_loss: 1.1515 - val_bounding_box_accuracy: 0.4940 -
val_class_label_accuracy: 0.6048
Epoch 25/50
62/62 [=====] - 5s 80ms/step - loss: 1.3154 -
bounding_box_loss: 0.0270 - class_label_loss: 1.2884 -
bounding_box_accuracy: 0.6904 - class_label_accuracy: 0.4806 -
val_loss: 1.2024 - val_bounding_box_loss: 0.0658 -
val_class_label_loss: 1.1366 - val_bounding_box_accuracy: 0.4456 -
val_class_label_accuracy: 0.6129
Epoch 26/50
62/62 [=====] - 4s 71ms/step - loss: 1.2792 -
bounding_box_loss: 0.0261 - class_label_loss: 1.2531 -
bounding_box_accuracy: 0.7100 - class_label_accuracy: 0.4952 -
val_loss: 1.2330 - val_bounding_box_loss: 0.0657 -
val_class_label_loss: 1.1673 - val_bounding_box_accuracy: 0.5343 -
val_class_label_accuracy: 0.5645
Epoch 27/50
62/62 [=====] - 4s 67ms/step - loss: 1.2521 -
bounding_box_loss: 0.0250 - class_label_loss: 1.2272 -
bounding_box_accuracy: 0.7060 - class_label_accuracy: 0.5184 -
val_loss: 1.2538 - val_bounding_box_loss: 0.0672 -
val_class_label_loss: 1.1866 - val_bounding_box_accuracy: 0.5262 -
val_class_label_accuracy: 0.6250
Epoch 28/50
62/62 [=====] - 4s 67ms/step - loss: 1.2841 -
bounding_box_loss: 0.0259 - class_label_loss: 1.2582 -
bounding_box_accuracy: 0.6944 - class_label_accuracy: 0.5119 -
val_loss: 1.2287 - val_bounding_box_loss: 0.0645 -
val_class_label_loss: 1.1641 - val_bounding_box_accuracy: 0.4637 -
val_class_label_accuracy: 0.5766
Epoch 29/50
62/62 [=====] - 4s 67ms/step - loss: 1.2416 -
bounding_box_loss: 0.0246 - class_label_loss: 1.2171 -
bounding_box_accuracy: 0.7176 - class_label_accuracy: 0.5275 -
val_loss: 1.2026 - val_bounding_box_loss: 0.0626 -
val_class_label_loss: 1.1401 - val_bounding_box_accuracy: 0.4435 -
val_class_label_accuracy: 0.6391
Epoch 30/50
62/62 [=====] - 4s 70ms/step - loss: 1.2011 -
bounding_box_loss: 0.0243 - class_label_loss: 1.1769 -
bounding_box_accuracy: 0.7368 - class_label_accuracy: 0.5265 -
val_loss: 1.2844 - val_bounding_box_loss: 0.0644 -
val_class_label_loss: 1.2200 - val_bounding_box_accuracy: 0.5000 -
val_class_label_accuracy: 0.5907
Epoch 31/50
62/62 [=====] - 5s 76ms/step - loss: 1.1748 -
bounding_box_loss: 0.0251 - class_label_loss: 1.1497 -

bounding_box_accuracy: 0.7171 - class_label_accuracy: 0.5562 -
val_loss: 1.1911 - val_bounding_box_loss: 0.0633 -
val_class_label_loss: 1.1278 - val_bounding_box_accuracy: 0.4254 -
val_class_label_accuracy: 0.6452
Epoch 32/50
62/62 [=====] - 5s 75ms/step - loss: 1.1858 -
bounding_box_loss: 0.0250 - class_label_loss: 1.1608 -
bounding_box_accuracy: 0.7141 - class_label_accuracy: 0.5356 -
val_loss: 1.1700 - val_bounding_box_loss: 0.0646 -
val_class_label_loss: 1.1054 - val_bounding_box_accuracy: 0.4516 -
val_class_label_accuracy: 0.6532
Epoch 33/50
62/62 [=====] - 5s 81ms/step - loss: 1.1598 -
bounding_box_loss: 0.0266 - class_label_loss: 1.1332 -
bounding_box_accuracy: 0.6566 - class_label_accuracy: 0.5451 -
val_loss: 1.2553 - val_bounding_box_loss: 0.0630 -
val_class_label_loss: 1.1923 - val_bounding_box_accuracy: 0.4315 -
val_class_label_accuracy: 0.6028
Epoch 34/50
62/62 [=====] - 4s 67ms/step - loss: 1.2429 -
bounding_box_loss: 0.0235 - class_label_loss: 1.2194 -
bounding_box_accuracy: 0.7171 - class_label_accuracy: 0.5224 -
val_loss: 1.2956 - val_bounding_box_loss: 0.0639 -
val_class_label_loss: 1.2317 - val_bounding_box_accuracy: 0.5101 -
val_class_label_accuracy: 0.5867
Epoch 35/50
62/62 [=====] - 5s 75ms/step - loss: 1.2917 -
bounding_box_loss: 0.0237 - class_label_loss: 1.2680 -
bounding_box_accuracy: 0.7252 - class_label_accuracy: 0.5209 -
val_loss: 1.2562 - val_bounding_box_loss: 0.0639 -
val_class_label_loss: 1.1923 - val_bounding_box_accuracy: 0.5302 -
val_class_label_accuracy: 0.5988
Epoch 36/50
62/62 [=====] - 5s 79ms/step - loss: 1.2452 -
bounding_box_loss: 0.0239 - class_label_loss: 1.2213 -
bounding_box_accuracy: 0.7211 - class_label_accuracy: 0.5280 -
val_loss: 1.2805 - val_bounding_box_loss: 0.0651 -
val_class_label_loss: 1.2153 - val_bounding_box_accuracy: 0.4617 -
val_class_label_accuracy: 0.6109
Epoch 37/50
62/62 [=====] - 5s 76ms/step - loss: 1.1452 -
bounding_box_loss: 0.0237 - class_label_loss: 1.1215 -
bounding_box_accuracy: 0.7463 - class_label_accuracy: 0.5638 -
val_loss: 1.2707 - val_bounding_box_loss: 0.0673 -
val_class_label_loss: 1.2034 - val_bounding_box_accuracy: 0.5081 -
val_class_label_accuracy: 0.5645
Epoch 38/50
62/62 [=====] - 4s 67ms/step - loss: 1.1639 -
bounding_box_loss: 0.0239 - class_label_loss: 1.1401 -

bounding_box_accuracy: 0.7514 - class_label_accuracy: 0.5618 -
val_loss: 1.2569 - val_bounding_box_loss: 0.0650 -
val_class_label_loss: 1.1919 - val_bounding_box_accuracy: 0.5302 -
val_class_label_accuracy: 0.6492
Epoch 39/50
62/62 [=====] - 5s 80ms/step - loss: 1.1810 -
bounding_box_loss: 0.0235 - class_label_loss: 1.1575 -
bounding_box_accuracy: 0.7297 - class_label_accuracy: 0.5628 -
val_loss: 1.2141 - val_bounding_box_loss: 0.0678 -
val_class_label_loss: 1.1463 - val_bounding_box_accuracy: 0.5202 -
val_class_label_accuracy: 0.5968
Epoch 40/50
62/62 [=====] - 5s 75ms/step - loss: 1.1938 -
bounding_box_loss: 0.0238 - class_label_loss: 1.1700 -
bounding_box_accuracy: 0.7327 - class_label_accuracy: 0.5456 -
val_loss: 1.2514 - val_bounding_box_loss: 0.0636 -
val_class_label_loss: 1.1878 - val_bounding_box_accuracy: 0.4899 -
val_class_label_accuracy: 0.6169
Epoch 41/50
62/62 [=====] - 5s 75ms/step - loss: 1.1973 -
bounding_box_loss: 0.0236 - class_label_loss: 1.1737 -
bounding_box_accuracy: 0.7201 - class_label_accuracy: 0.5466 -
val_loss: 1.2746 - val_bounding_box_loss: 0.0677 -
val_class_label_loss: 1.2069 - val_bounding_box_accuracy: 0.5544 -
val_class_label_accuracy: 0.5887
Epoch 42/50
62/62 [=====] - 5s 79ms/step - loss: 1.1986 -
bounding_box_loss: 0.0237 - class_label_loss: 1.1749 -
bounding_box_accuracy: 0.7302 - class_label_accuracy: 0.5517 -
val_loss: 1.2887 - val_bounding_box_loss: 0.0648 -
val_class_label_loss: 1.2239 - val_bounding_box_accuracy: 0.5141 -
val_class_label_accuracy: 0.5948
Epoch 43/50
62/62 [=====] - 4s 67ms/step - loss: 1.1955 -
bounding_box_loss: 0.0231 - class_label_loss: 1.1724 -
bounding_box_accuracy: 0.7458 - class_label_accuracy: 0.5557 -
val_loss: 1.3193 - val_bounding_box_loss: 0.0652 -
val_class_label_loss: 1.2541 - val_bounding_box_accuracy: 0.4980 -
val_class_label_accuracy: 0.6089
Epoch 44/50
62/62 [=====] - 5s 75ms/step - loss: 1.1940 -
bounding_box_loss: 0.0230 - class_label_loss: 1.1710 -
bounding_box_accuracy: 0.7458 - class_label_accuracy: 0.5547 -
val_loss: 1.2351 - val_bounding_box_loss: 0.0663 -
val_class_label_loss: 1.1687 - val_bounding_box_accuracy: 0.5343 -
val_class_label_accuracy: 0.6129
Epoch 45/50
62/62 [=====] - 4s 70ms/step - loss: 1.2131 -
bounding_box_loss: 0.0225 - class_label_loss: 1.1906 -

```

bounding_box_accuracy: 0.7312 - class_label_accuracy: 0.5456 -
val_loss: 1.2993 - val_bounding_box_loss: 0.0670 -
val_class_label_loss: 1.2323 - val_bounding_box_accuracy: 0.4798 -
val_class_label_accuracy: 0.6290
Epoch 46/50
62/62 [=====] - 5s 76ms/step - loss: 1.1978 -
bounding_box_loss: 0.0230 - class_label_loss: 1.1748 -
bounding_box_accuracy: 0.7393 - class_label_accuracy: 0.5547 -
val_loss: 1.3155 - val_bounding_box_loss: 0.0652 -
val_class_label_loss: 1.2503 - val_bounding_box_accuracy: 0.5020 -
val_class_label_accuracy: 0.5685
Epoch 47/50
62/62 [=====] - 4s 68ms/step - loss: 1.1548 -
bounding_box_loss: 0.0227 - class_label_loss: 1.1321 -
bounding_box_accuracy: 0.7438 - class_label_accuracy: 0.5688 -
val_loss: 1.3051 - val_bounding_box_loss: 0.0640 -
val_class_label_loss: 1.2410 - val_bounding_box_accuracy: 0.5060 -
val_class_label_accuracy: 0.6048
Epoch 48/50
62/62 [=====] - 5s 81ms/step - loss: 1.1239 -
bounding_box_loss: 0.0228 - class_label_loss: 1.1011 -
bounding_box_accuracy: 0.7292 - class_label_accuracy: 0.5875 -
val_loss: 1.2697 - val_bounding_box_loss: 0.0659 -
val_class_label_loss: 1.2038 - val_bounding_box_accuracy: 0.4919 -
val_class_label_accuracy: 0.6310
Epoch 49/50
62/62 [=====] - 4s 67ms/step - loss: 1.1470 -
bounding_box_loss: 0.0235 - class_label_loss: 1.1235 -
bounding_box_accuracy: 0.7211 - class_label_accuracy: 0.5769 -
val_loss: 1.2435 - val_bounding_box_loss: 0.0637 -
val_class_label_loss: 1.1799 - val_bounding_box_accuracy: 0.4516 -
val_class_label_accuracy: 0.6190
Epoch 50/50
62/62 [=====] - 4s 66ms/step - loss: 1.1293 -
bounding_box_loss: 0.0224 - class_label_loss: 1.1069 -
bounding_box_accuracy: 0.7156 - class_label_accuracy: 0.5749 -
val_loss: 1.2575 - val_bounding_box_loss: 0.0653 -
val_class_label_loss: 1.1922 - val_bounding_box_accuracy: 0.5262 -
val_class_label_accuracy: 0.5968

```

Visualizes our three loss components: the class label loss, bounding box loss, and total loss

```

# plot the total loss, label loss, and bounding box loss
epochs=50
lossNames = ["loss", "class_label_loss", "bounding_box_loss"]
N = np.arange(0, epochs)
plt.style.use("ggplot")
(fig, ax) = plt.subplots(3, 1, figsize=(13, 13))


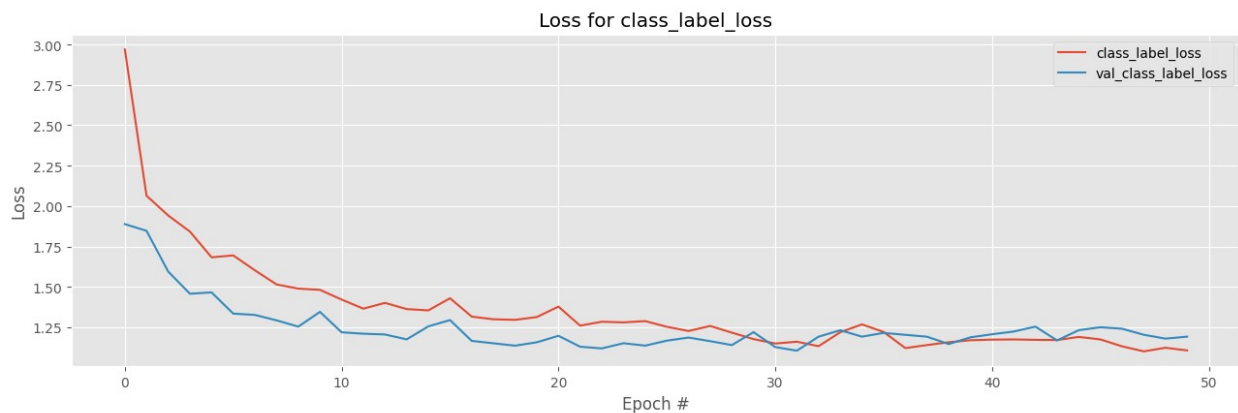

```



```

# loop over the loss names
for (i, l) in enumerate(lossNames):
    # plot the loss for both the training and validation data
    title = "Loss for {}".format(l) if l != "loss" else "Total loss"
    ax[i].set_title(title)
    ax[i].set_xlabel("Epoch #")
    ax[i].set_ylabel("Loss")
    ax[i].plot(N, H.history[l], label=l)
    ax[i].plot(N, H.history["val_" + l], label="val_" + l)
    ax[i].legend()
# save the losses figure and create a new figure for the accuracies
plt.tight_layout()


```



Visualizes two accuracy components: bounding box loss, class label loss

```
# create a new figure for the accuracies
plt.style.use("ggplot")
plt.figure()
plt.plot(N, H.history["bounding_box_accuracy"],
         label="bounding_box_accuracy")
plt.plot(N, H.history["val_bounding_box_accuracy"],
         label="val_bounding_box_accuracy")
plt.title("Bounding Box Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Accuracy")
plt.legend(loc="lower left")


<matplotlib.legend.Legend at 0x7fcf08075b70>
```



```
# create a new figure for the accuracies
plt.style.use("ggplot")
plt.figure()
plt.plot(N, H.history["class_label_accuracy"],
         label="class_label_train_acc")
plt.plot(N, H.history["val_class_label_accuracy"],
         label="val_class_label_acc")
```

```
plt.title("Class Label Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Accuracy")
plt.legend(loc="lower left")

<matplotlib.legend.Legend at 0x7fcf080d26b0>
```



Predict multi-class objects

```
img =
'/content/crop-disease-ghana/input/Corn/Corn__Cercospora_Leaf_Spot/
images/20230524_104642.jpg'
#img =
'/content/crop-disease-ghana/input/Tomato/Tomato__Septoria/images/
2N8A0600.JPG'

# loop over the images that we'll be testing using our bounding box
# regression model
#for imagePath in imagePaths:
    # load the input image (in Keras format) from disk and preprocess
    # it, scaling the pixel intensities to the range [0, 1]

image = load_img(img, target_size=(224, 224))
```

```

image = img_to_array(image) / 255.0
image = np.expand_dims(image, axis=0)
# predict the bounding box of the object along with the class
# label
(boxPreds, labelPreds) = model.predict(image)
(startX, startY, endX, endY) = boxPreds[0]
# determine the class label with the largest predicted
# probability
i = np.argmax(labelPreds, axis=1)
label = label_binarizer.classes_[i][0]

1/1 [=====] - 0s 487ms/step

print(startX, startY, endX, endY)

0.54288334 0.51715934 0.6258746 0.61109596

print(label)

Corn Common Rust

# load the input image (in OpenCV format), resize it such that it
# fits on our screen, and grab its dimensions
from imutils import paths

image = cv2.imread(img)
image = image.copy()
image = cv2.resize(image, (400, 600))
#image = imutils.resize(image, width=600)
(h, w) = image.shape[:2]
# scale the predicted bounding box coordinates based on the image
# dimensions
startX = int(startX * w)
startY = int(startY * h)
endX = int(endX * w)
endY = int(endY * h)
print(startX, startY, endX, endY)
# draw the predicted bounding box and class label on the image
y = startY - 10 if startY - 10 > 10 else startY + 10
cv2.putText(image, label, (startX, y), cv2.FONT_HERSHEY_SIMPLEX, 0.65,
(0, 255, 0), 2)
cv2.rectangle(image, (startX, startY), (endX, endY), (0, 255, 0), 2)
# show the output image
cv2.imshow(image)
cv2.waitKey(0)

217 310 250 366

```



Model Prediction and Matrix

```
scores = model.evaluate(testImages, testTargets)
print(f"Test Bounding Box Accuracy: {scores[3]*100}")
print(f"Test Class Label Accuracy: {scores[4]*100}")

16/16 [=====] - 1s 48ms/step - loss: 1.2575 -
bounding_box_loss: 0.0653 - class_label_loss: 1.1922 -
bounding_box_accuracy: 0.5262 - class_label_accuracy: 0.5968
Test Bounding Box Accuracy: 52.620965242385864
Test Class Label Accuracy: 59.67742204666138

pred_test=model.predict(testImages)

16/16 [=====] - 1s 42ms/step

# Reshape to remove the redundant dimension
converted_coordinates = testBBboxes.reshape((testBBboxes.shape[0],
testBBboxes.shape[1]))

print(converted_coordinates)

[[0.45542297 0.66766346 0.5261052 0.70759094]
 [0.4472397 0.45805028 0.5366876 0.5310316 ]
 [0.5735199 0.22531866 0.67679214 0.2906433 ]
 ...
 [0.441478 0.43594 0.4965056 0.5346279 ]
 [0.6053255 0.32573956 0.64561933 0.36646754]
 [0.30718216 0.3675907 0.39954984 0.44447142]]

from sklearn.metrics import mean_squared_error

mse_bbox = mean_squared_error(converted_coordinates, pred_test[0])
```

MSE of bounding box predictions

```
print(f"Test Bounding Box Mean Square Error: {mse_bbox*100}")

Test Bounding Box Mean Square Error: 6.52831494808197

binary_predictions = np.argmax(pred_test[1] , axis=1).astype(int)

# threshold = 0.5
# binary_predictions = (pred_test[1] > threshold).astype(int)

testLabelsNonBinary=testLabels.argmax(axis=1)
```

F1 Score, Accuracy, Precision and Recall of class label prediction

```
# accuracy: (tp + tn) / (p + n)
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
```

```
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

f1_test = f1_score(testLabelsNonBinary, binary_predictions,
average='weighted')
print(f"Test Classifications F1 score: {f1_test*100}")
accuracy = accuracy_score(testLabelsNonBinary, binary_predictions)
print(f"Test Classifications accracy: {accuracy*100}")
# precision tp / (tp + fp)
precision = precision_score(testLabelsNonBinary, binary_predictions,
average='weighted')
print(f"Test Classifications precision: {precision*100}")
# recall: tp / (tp + fn)
recall = recall_score(testLabelsNonBinary, binary_predictions,
average='weighted')
print(f"Test Classifications recall: {recall*100}")

Test Classifications F1 score: 59.705838466327286
Test Classifications accracy: 59.67741935483871
Test Classifications precision: 63.84565384949366
Test Classifications recall: 59.67741935483871
```